

# HW02 — ChBE 413

**Name:** Mikkel Flores

**Due Date:** September 11, 2025

## Question 1: Processing and Visualizing Molecular Datasets with Pandas

```
In [185... import rdkit
from rdkit import Chem
from rdkit.Chem import AllChem
import pandas as pd
```

### Part a

```
In [186... df = pd.read_csv("Tm_200_subset_raw.csv")
df.head()
```

```
Out[186...
```

|   | smiles   | Tm     | num_atms | dipole   | quadrupole |
|---|--|--------|----------|----------|------------|
| 0 | <chem>CC1=C(C=CC(O)=N1)[N+](=[O-])=O</chem>        | 508.15 | 17.0     | 4.473978 | 23.191697  |
| 1 | <chem>COC1=C(N)C=C(C=C1)C(=O)N(C)C</chem>          | 393.15 | 28.0     | 4.772571 | 21.767880  |
| 2 | <chem>CC1=CC=C(Cl)C(N)=C1</chem>                   | 303.95 | 17.0     | 2.216088 | 7.920099   |
| 3 | <chem>BrC1(C(=O)C2=CC=CC=C2C1=O)C1=CC=CC=C1</chem> | 379.15 | 27.0     | 3.919235 | 21.700726  |
| 4 | <chem>NC1=C(Cl)C=C(Cl)C=C1I</chem>                 | 353.15 | 14.0     | 2.157331 | 4.988358   |

```
In [187... duplicates = df.duplicated().sum()
print(f"Number of duplicate rows: {duplicates}")
```

Number of duplicate rows: 12

In [188... `df.drop_duplicates()`

Out [188...

|     | smiles   | Tm     | num_atms | dipole   | quadrupole |
|-----|--|--------|----------|----------|------------|
| 0   | <chem>CC1=C(C=CC(O)=N1)[N+](O-)=O</chem>           | 508.15 | 17.0     | 4.473978 | 23.191697  |
| 1   | <chem>COC1=C(N)C=C(C=C1)C(=O)N(C)C</chem>          | 393.15 | 28.0     | 4.772571 | 21.767880  |
| 2   | <chem>CC1=CC=C(Cl)C(N)=C1</chem>                   | 303.95 | 17.0     | 2.216088 | 7.920099   |
| 3   | <chem>BrC1(C(=O)C2=CC=CC=C2C1=O)C1=CC=CC=C1</chem> | 379.15 | 27.0     | 3.919235 | 21.700726  |
| 4   | <chem>NC1=C(Cl)C=C(Cl)C=C1I</chem>                 | 353.15 | 14.0     | 2.157331 | 4.988358   |
| ... | ...  | ...    | ...      | ...      | ...        |
| 208 | <chem>O=CC1=CN=C(S1)N1CCOCC1</chem>                | 437.15 | 23.0     | 4.874555 | 18.470379  |
| 209 | <chem>CC(C)N(C)S(=O)(=O)C1=CC(N)=C(C)C=C1</chem>   | 391.15 | 34.0     | NaN      | 21.523472  |
| 210 | <chem>NC1=C(N=CN=C1)N1CCCC1</chem>                 | 430.15 | 2400.0   | 3.889375 | 11.228281  |
| 211 | <chem>CC(=O)N[C@@H](CC(O)=O)C(N)=O</chem>          | 433.15 | 22.0     | 3.437255 | 11.903019  |
| 212 | <chem>CN(C)C(=O)C1=CC(=CC=C1)S(N)(=O)=O</chem>     | 395.15 | 27.0     | 4.158610 | 19.023166  |

201 rows × 5 columns

### Part b

In [189... `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 213 entries, 0 to 212
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  -
0   smiles      213 non-null    object
1   Tm           213 non-null    float64
2   num_atms     210 non-null    float64
3   dipole       205 non-null    float64
4   quadrupole   213 non-null    float64
dtypes: float64(4), object(1)
memory usage: 8.4+ KB
```

```
In [190... df.isnull().sum()
```

```
Out[190... smiles      0
Tm           0
num_atms     3
dipole       8
quadrupole   0
dtype: int64
```

There are 3 entries missing in the num\_atms column and 8 entries missing in the dipole column.

### Part c

```
In [191... mean_col_vals = df[['Tm', 'num_atms', 'dipole', 'quadrupole']].mean()
print("The mean column values:",
      "\n", mean_col_vals)
```

```
The mean column values:
Tm           419.614930
num_atms     38.090476
dipole       23.218965
quadrupole   301.074375
dtype: float64
```

```
In [192... # Checking skewness and percentage of missing values
df_skew = df[['Tm', 'num_atms', 'dipole', 'quadrupole']].skew()
df_missing = (df[['Tm', 'num_atms', 'dipole', 'quadrupole']].isnull().mean()*100).round(2)
print("Skewness of the columns:"
      , "\n", df_skew)
```

```
, "\n\nPercentage of missing values in the columns:"
, "\n", df_missing)
```

Skewness of the columns:

```
Tm          11.555019
num_atms    14.412902
dipole      14.316706
quadrupole  14.594364
dtype: float64
```

Percentage of missing values in the columns:

```
Tm          0.00
num_atms    1.41
dipole      3.76
quadrupole  0.00
dtype: float64
```

Mean imputation is not necessarily a good approximation here because the columns show high positive skew, which demonstrates that the mean is not representative and may bias results. The article notes that mean imputation does not account for correlations between features, so approximating with the column means ignores those relationships, subsequently worsening models or further multivariate analyses. Even though the percentage of missing values is small, imputation with the mean reduces the variance and can artificially make a model more confident since uncertainty is not accounted for these imputations.

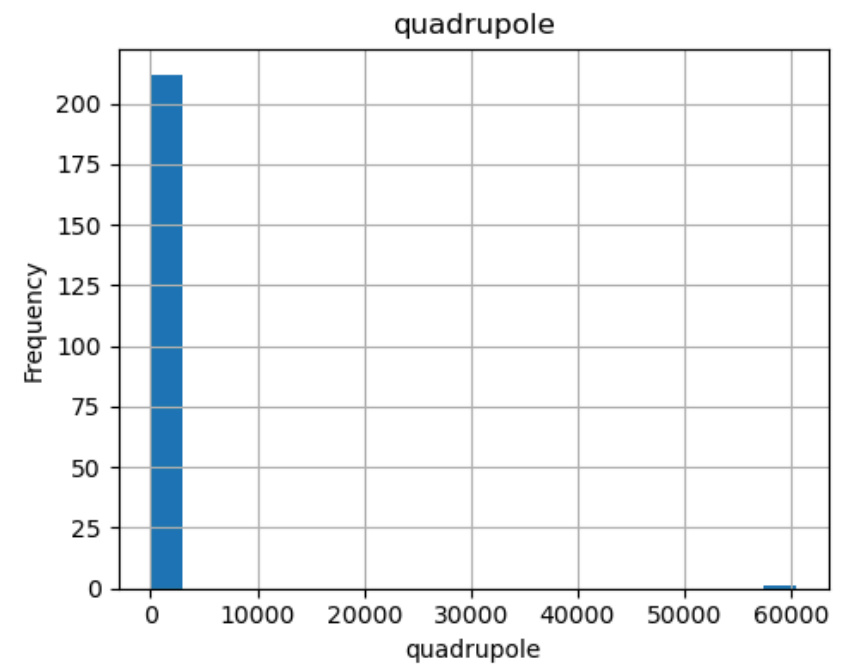
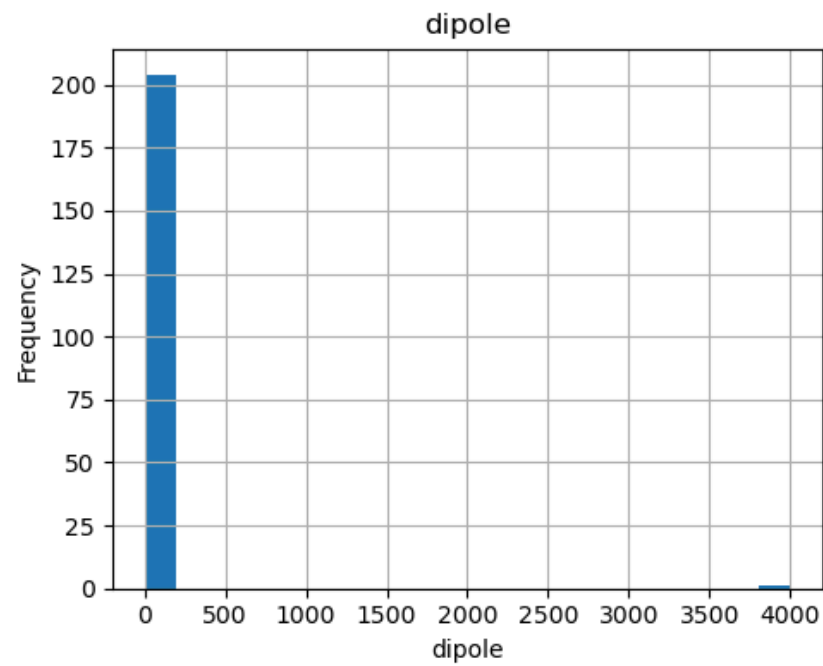
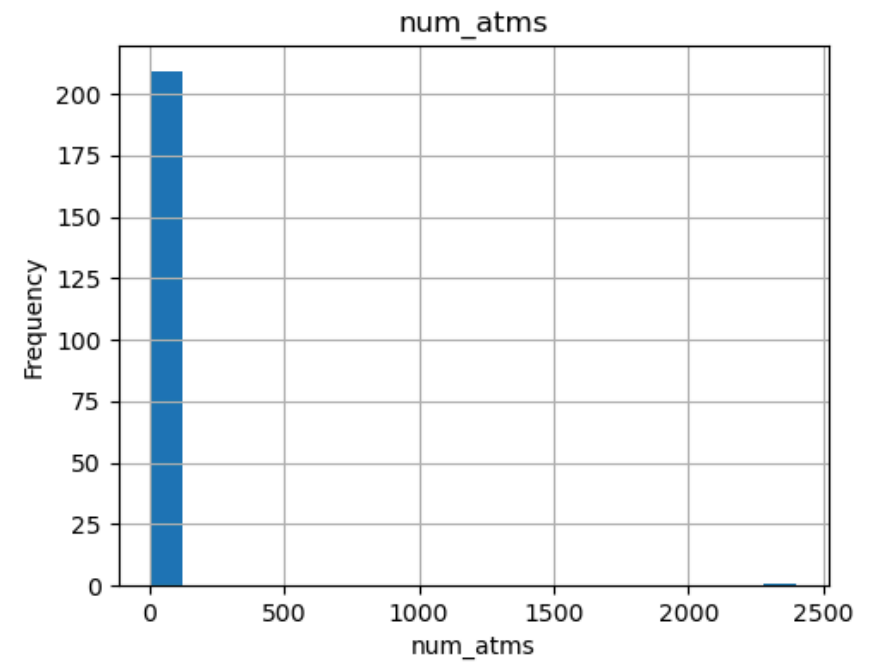
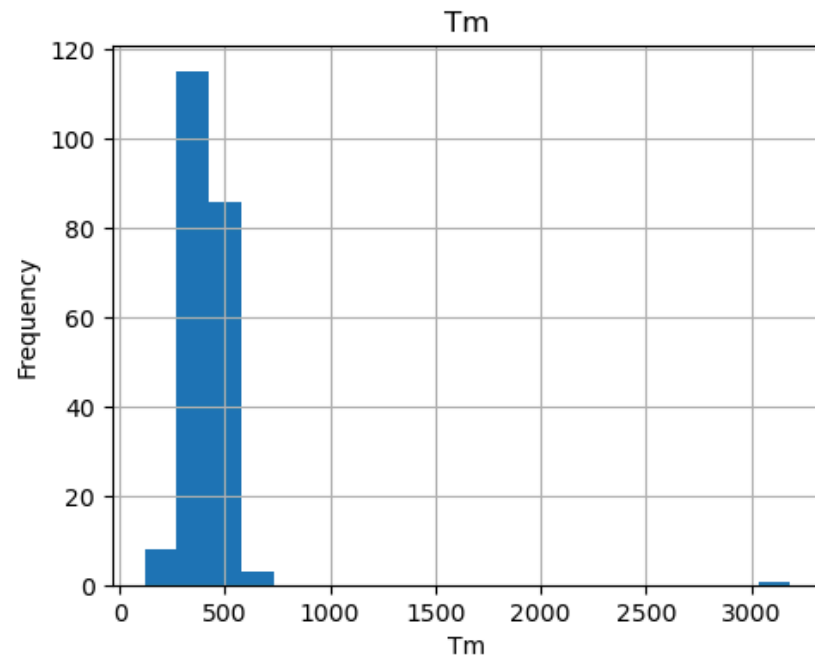
## Part d

```
In [193]: import matplotlib.pyplot as plt
cols = ['Tm', 'num_atms', 'dipole', 'quadrupole']
num_cols = len(cols)
num_plots_per_col = 2
num_plots_per_row = (num_cols + num_plots_per_col - 1) // num_plots_per_col

fig, axes = plt.subplots(num_plots_per_row, num_plots_per_col, figsize=(10, 4*num_plots_per_row))
axes = axes.flatten()

for i, col in enumerate(cols):
    ax = axes[i]
    df[col].hist(bins=20, ax=ax)
    ax.set_title(col)
    ax.set_xlabel(col)
```

```
ax.set_ylabel("Frequency")  
ax.grid(True)  
  
for j in range(i+1, len(axes)):  
    axes[j].axis('off')  
  
plt.tight_layout()  
plt.show()
```



```
In [194... description = df[cols].describe().transpose()
print(description.to_string())
```

|            | count | mean       | std         | min        | 25%        | 50%        | 75%        | max          |
|------------|-------|------------|-------------|------------|------------|------------|------------|--------------|
| Tm         | 213.0 | 419.614930 | 205.510776  | 119.150000 | 361.150000 | 413.150000 | 451.150000 | 3186.500000  |
| num_atms   | 210.0 | 38.090476  | 164.061916  | 5.000000   | 20.000000  | 25.500000  | 31.750000  | 2400.000000  |
| dipole     | 205.0 | 23.218965  | 279.247777  | 0.000000   | 2.371265   | 3.486184   | 5.003426   | 4001.832900  |
| quadrupole | 213.0 | 301.074375 | 4141.268430 | 0.366284   | 9.455462   | 15.351293  | 22.387872  | 60456.931093 |

```
In [195... means = description['mean']
stds = description['std']
outlier_mult = 3
flag_mask = pd.Series(False, index=df.index)

for c in cols:
    lower_bound = means[c] - outlier_mult * stds[c]
    upper_bound = means[c] + outlier_mult * stds[c]
    outliers = (df[c] < lower_bound) | (df[c] > upper_bound)
    print(f"{c}: flagged {outliers.sum()} rows (bounds {lower_bound:.2f} .. {upper_bound:.2f})")
    flag_mask |= outliers

flagged = df.loc[flag_mask, ['smiles'] + cols]
print("\nFlagged rows:", len(flagged))
display(flagged)
print("\nIndices to drop:", flagged.index.tolist())
```

Tm: flagged 1 rows (bounds -196.92 .. 1036.15)  
 num\_atms: flagged 1 rows (bounds -454.10 .. 530.28)  
 dipole: flagged 1 rows (bounds -814.52 .. 860.96)  
 quadrupole: flagged 1 rows (bounds -12122.73 .. 12724.88)

Flagged rows: 4

|            | smiles  | Tm      | num_atms | dipole      | quadrupole   |
|------------|---|---------|----------|-------------|--------------|
| <b>13</b>  | O=CC1=CC=CC=C1N1CCOCC1                        | 3186.50 | 27.0     | 3.318658    | 15.881392    |
| <b>26</b>  | CS(=O)(=O)C1=CC=C(C=C1)C(O)C(CO)NC(=O)C(Cl)Cl | 438.15  | 36.0     | 4001.832900 | 27.426057    |
| <b>191</b> | COC(=O)C1=CC=C(I)C=C1O                        | 342.15  | 19.0     | 2.664381    | 60456.931093 |
| <b>210</b> | NC1=C(N=CN=C1)N1CCCC1                         | 430.15  | 2400.0   | 3.889375    | 11.228281    |

Indices to drop: [13, 26, 191, 210]

```
In [196... display(df.loc[flagged.index, ['smiles'] + cols])
print("\nSMILES and indices of dropped rows:")
print("SMILES list:", df.loc[flagged.index, 'smiles'].tolist())
print("Indices:", flagged.index.tolist())
```

|            | smiles   | Tm      | num_atms | dipole      | quadrupole   |
|------------|--|---------|----------|-------------|--------------|
| <b>13</b>  | <chem>O=CC1=CC=CC=C1N1CCOCC1</chem>                        | 3186.50 | 27.0     | 3.318658    | 15.881392    |
| <b>26</b>  | <chem>CS(=O)(=O)C1=CC=C(C=C1)C(O)C(CO)NC(=O)C(Cl)Cl</chem> | 438.15  | 36.0     | 4001.832900 | 27.426057    |
| <b>191</b> | <chem>COC(=O)C1=CC=C(I)C=C1O</chem>                        | 342.15  | 19.0     | 2.664381    | 60456.931093 |
| <b>210</b> | <chem>NC1=C(N=CN=C1)N1CCCC1</chem>                         | 430.15  | 2400.0   | 3.889375    | 11.228281    |

SMILES and indices of dropped rows:

SMILES list: ['O=CC1=CC=CC=C1N1CCOCC1', 'CS(=O)(=O)C1=CC=C(C=C1)C(O)C(CO)NC(=O)C(Cl)Cl', 'COC(=O)C1=CC=C(I)C=C1O', 'NC1=C(N=CN=C1)N1CCCC1']

Indices: [13, 26, 191, 210]

```
In [197... df.drop(index=flagged.index, inplace=True)
df.reset_index(drop=True, inplace=True)
df.to_csv("Tm_200_subset_cleaned.csv", index=False)
print("rows after cleaning:", len(df))
```

rows after cleaning: 209

### Part e

```
In [198... df_clean = pd.read_csv("Tm_200_subset_cleaned.csv")
df_clean.drop_duplicates(inplace=True)
df_clean.head()
df_clean.info()
print("\nNaN Values:", "\n", df_clean.isnull().sum())
print("\nDuplicated Values:", df_clean.duplicated().sum())
```



```
<class 'pandas.core.frame.DataFrame'>
Index: 197 entries, 0 to 208
Data columns (total 5 columns):
#   Column      Non-Null Count  Dtype
---  -
0    smiles      197 non-null    object
1    Tm           197 non-null    float64
2    num_atms     194 non-null    float64
3    dipole       189 non-null    float64
4    quadrupole   197 non-null    float64
dtypes: float64(4), object(1)
memory usage: 9.2+ KB
```

NaN Values:

```
smiles      0
Tm           0
num_atms     3
dipole       8
quadrupole   0
dtype: int64
```

Duplicated Values: 0

```
In [199... cols = ['Tm', 'num_atms', 'dipole', 'quadrupole']

corr = df_clean[cols].corr(method='pearson')
print("Pearson correlation matrix:\n", corr)
```

Pearson correlation matrix:

|            |          | Tm       | num_atms | dipole   | quadrupole |
|------------|----------|----------|----------|----------|------------|
| Tm         | 1.000000 | 0.107806 | 0.332250 | 0.358435 |            |
| num_atms   | 0.107806 | 1.000000 | 0.105737 | 0.367393 |            |
| dipole     | 0.332250 | 0.105737 | 1.000000 | 0.351995 |            |
| quadrupole | 0.358435 | 0.367393 | 0.351995 | 1.000000 |            |

From the Pearson correlation matrix, it seems that the two columns most correlated to each other are num\_atms and quadrupole with an r value of 0.367. The least correlated pair seems to be num\_atms and dipole ( $r = 0.106$ ). All pairwise Pearson r values are positive, indicating there are no anti-correlations between any of the columns.

```
In [200... # Showing pairwise counts of non-NaN values
print("\nPairwise counts of non-NaN values:")
```

```
for i in cols:
    for j in cols:
        print(f"{i:11s} vs {j:11s}: {df_clean[[i, j]].dropna().shape[0]}")
```

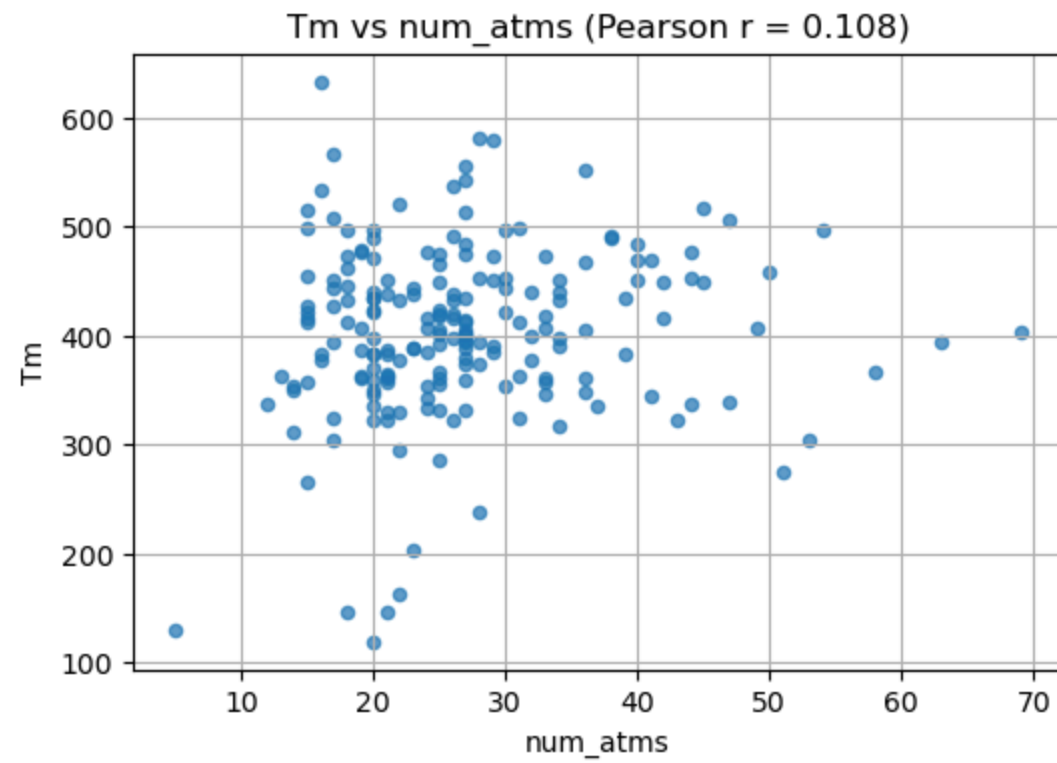
Pairwise counts of non-NaN values:

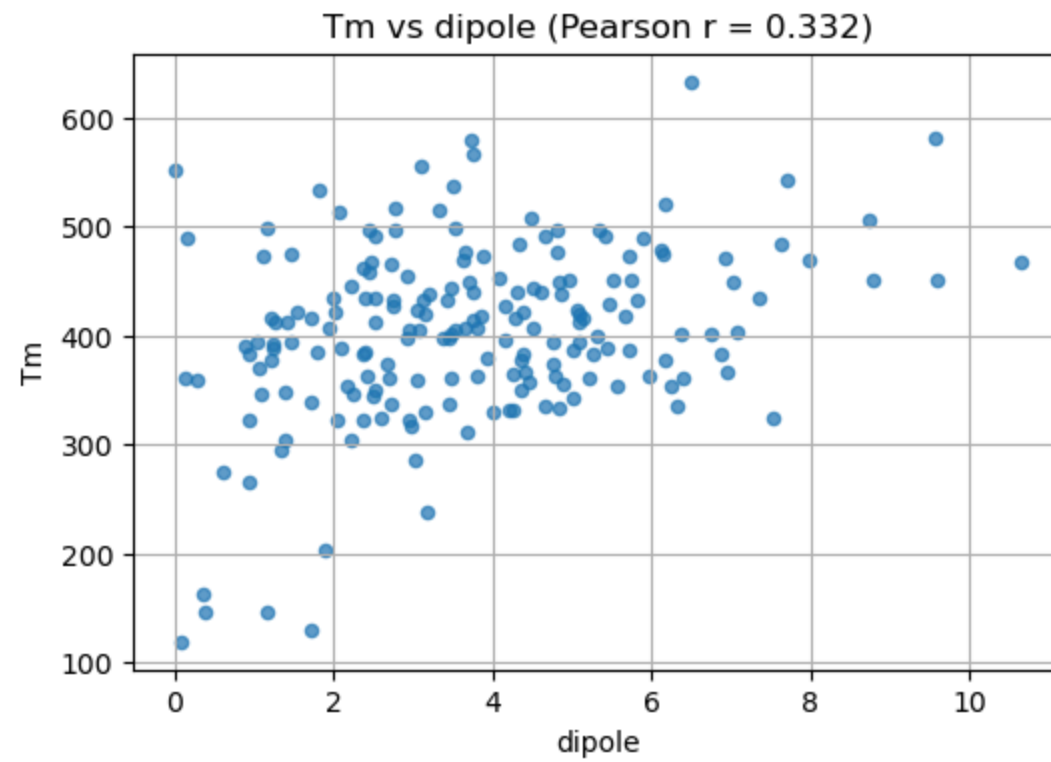
```
Tm          vs Tm          : 197
Tm          vs num_atms    : 194
Tm          vs dipole      : 189
Tm          vs quadrupole  : 197
num_atms     vs Tm          : 194
num_atms     vs num_atms    : 194
num_atms     vs dipole      : 186
num_atms     vs quadrupole  : 194
dipole       vs Tm          : 189
dipole       vs num_atms    : 186
dipole       vs dipole      : 189
dipole       vs quadrupole  : 189
quadrupole   vs Tm          : 197
quadrupole   vs num_atms    : 194
quadrupole   vs dipole      : 189
quadrupole   vs quadrupole  : 197
```

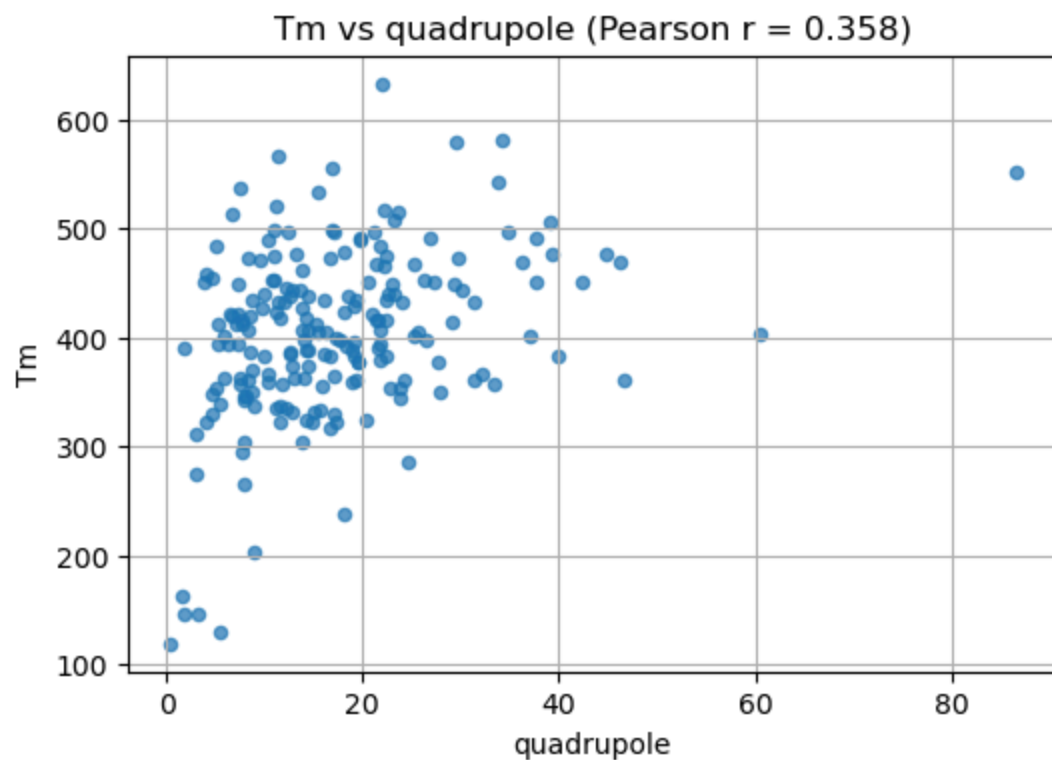
### Part f

```
In [201... for feat in ['num_atms', 'dipole', 'quadrupole']:
    data = df_clean[['Tm', feat]].dropna()
    x = data[feat]
    y = data['Tm']
    r = x.corr(y)

    plt.figure(figsize=(6, 4))
    plt.scatter(x, y, alpha=0.7, s=20)
    plt.title(f'Tm vs {feat} (Pearson r = {r:.3f})')
    plt.xlabel(feat)
    plt.ylabel('Tm')
    plt.grid(True)
    plt.show()
```







From the plots, dipole and quadrupole exhibits the strongest visual correlation with Tm, which is consistent with the computed correlation coefficients in e.

### Part g

```
In [202... lowest3 = df_clean.nsmallest(3, 'Tm').copy()
highest3 = df_clean.nlargest(3, 'Tm').copy()

lowest3_smiles = lowest3['smiles'].tolist()
highest3_smiles = highest3['smiles'].tolist()
print("Lowest 3 Tm SMILES:", lowest3_smiles)
print("Highest 3 Tm SMILES:", highest3_smiles)

print("\nLowest 3 Tm molecules:")
display(lowest3[['smiles'] + cols])
```

```
print("\nHighest 3 Tm molecules:")
display(highest3[['smiles'] + cols])
```

Lowest 3 Tm SMILES: ['CCCC(C)C', 'CF', 'CCC(CC)C=C']

Highest 3 Tm SMILES: ['NC1=NC2=C(N=CN2)C(=O)N1', 'NC1=C(C#N)C(=O)C2CC3=C(CN12)C=CC=C3', 'OC(=O)C1=CC(NC(=O)C2=CC=C(Br)C=C2)=CC=C1']

Lowest 3 Tm molecules:

|           | smiles     | Tm     | num_atms | dipole   | quadrupole |
|-----------|------------|--------|----------|----------|------------|
| <b>12</b> | CCCC(C)C   | 119.15 | 20.0     | 0.067055 | 0.366284   |
| <b>53</b> | CF         | 130.15 | 5.0      | 1.700645 | 5.380003   |
| <b>85</b> | CCC(CC)C=C | 146.15 | 21.0     | 0.386677 | 1.692360   |

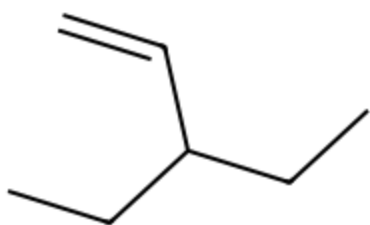
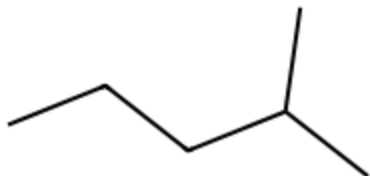
Highest 3 Tm molecules:

|            | smiles                                   | Tm     | num_atms | dipole   | quadrupole |
|------------|--|--------|----------|----------|------------|
| <b>80</b>  | NC1=NC2=C(N=CN2)C(=O)N1                  | 633.15 | 16.0     | 6.488483 | 21.894593  |
| <b>9</b>   | NC1=C(C#N)C(=O)C2CC3=C(CN12)C=CC=C3      | 581.15 | 28.0     | 9.580572 | 34.144607  |
| <b>197</b> | OC(=O)C1=CC(NC(=O)C2=CC=C(Br)C=C2)=CC=C1 | 580.15 | 29.0     | 3.733080 | 29.455901  |

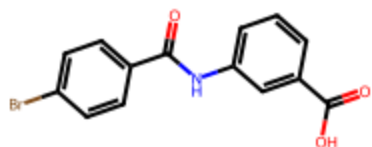
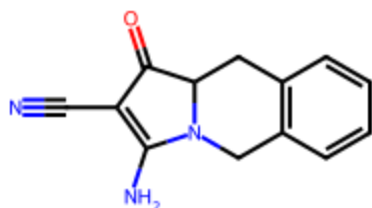
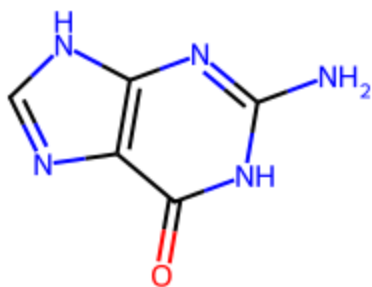
```
In [203... print("\n2D depictions (lowest 3):")
for sm in lowest3_smiles:
    mol = Chem.MolFromSmiles(sm)
    img = Chem.Draw.MolToImage(mol, size=(200, 200))
    display(img)

print("\n2D depictions (highest 3):")
for sm in highest3_smiles:
    mol = Chem.MolFromSmiles(sm)
    img = Chem.Draw.MolToImage(mol, size=(200, 200))
    display(img)
```

2D depictions (lowest 3):



2D depictions (highest 3):



```
In [204... # Exploring extra functionalities from the rdkit.Chem module since this is new to me; I thought it was cool
from rdkit.Chem import Descriptors, rdMolDescriptors, rdchem

rows = pd.concat([lowest3, highest3], ignore_index=True)
```



```
descrs = []
for _, r in rows.iterrows():
    sm = r['smiles']
    m = Chem.MolFromSmiles(sm)
    if m is None:
        descrs.append({
            'smiles': sm, 'Tm': r['Tm'],
            'MolWt': None, 'HeavyAtoms': None,
            'NumRings': None, 'NumRotatable': None,
            'NumHDonors': None, 'NumHAcceptors': None,
            'BondSingle': None, 'BondDouble': None, 'BondTriple': None, 'BondAromatic': None
        })
        continue

    # bond counts
    single = 0
    double = 0
    triple = 0
    aromatic = 0
    for b in m.GetBonds():
        bt = b.GetBondType()
        if bt == rdchem.BondType.SINGLE:
            single += 1
        elif bt == rdchem.BondType.DOUBLE:
            double += 1
        elif bt == rdchem.BondType.TRIPLE:
            triple += 1
    # aromatic detection
    if b.GetIsAromatic():
        aromatic += 1

    descrs.append({
        'smiles': sm,
        'Tm': float(r['Tm']),
        'MolWt': float(Descriptors.MolWt(m)),
        'HeavyAtoms': int(m.GetNumHeavyAtoms()),
        'NumRings': int(rdMolDescriptors.CalcNumRings(m)),
        'NumRotatable': int(rdMolDescriptors.CalcNumRotatableBonds(m)),
        'NumHDonors': int(Descriptors.NumHDonors(m)),
        'NumHAcceptors': int(Descriptors.NumHAcceptors(m)),
        'BondSingle': int(single),
        'BondDouble': int(double),
```

```

        'BondTriple': int(triple),
        'BondAromatic': int(aromatic),
    })

descr_df = pd.DataFrame(descrs)
print("Descriptor table for lowest3 then highest3 (includes bond counts):")
display(descr_df)

```

Descriptor table for lowest3 then highest3 (includes bond counts):

|   | smiles                                   | Tm     | MolWt   | HeavyAtoms | NumRings | NumRotatable | NumHDonors |
|---|--|--------|---------|------------|----------|--------------|------------|
| 0 | CCCC(C)C                                 | 119.15 | 86.178  | 6          | 0        | 2            | 0          |
| 1 | CF                                       | 130.15 | 34.033  | 2          | 0        | 0            | 0          |
| 2 | CCC(CC)C=C                               | 146.15 | 98.189  | 7          | 0        | 3            | 0          |
| 3 | NC1=NC2=C(N=CN2)C(=O)N1                  | 633.15 | 151.129 | 11         | 2        | 0            | 3          |
| 4 | NC1=C(C#N)C(=O)C2CC3=C(CN12)C=CC=C3      | 581.15 | 225.251 | 17         | 3        | 0            | 1          |
| 5 | OC(=O)C1=CC(NC(=O)C2=CC=C(Br)C=C2)=CC=C1 | 580.15 | 320.142 | 19         | 2        | 3            | 2          |

While I included multiple qualitative molecular descriptors from rd.Chem, it seems that molecular size (MolWt and HeavyAtoms) and intermolecular interactions (dipole, polarity, NumHDonors/NumHAcceptors) seem to most correlate with increasing Tm. The larger molecules have greater polarizability and more potential for London dispersion forces, both of which scale with the increased surface area from the presence of Heavy Atoms. This cohesive energy creates a crystal lattice that is more "tightly packed" which subsequently increases Tm relative to the smaller molecules. Going off of this, the number of polar functional groups and H-bond donors/acceptors also correlate with Tm. These atoms create strong dipole-dipole and hydrogen bonds that increase the cohesion energy which also raise Tm when compared to molecules with less of these types of atoms.